

Importing necessary libraries

```
In [1]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
get_ipython().run_line_magic('matplotlib', 'inline')

from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn import metrics

import seaborn as sns
```

```
In [2]: df=pd.read_csv('Iris.csv') #reading the iris dataset which is in .CSV format
```

```
In [4]: df.head(10)
```

```
Out[4]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
5	6	5.4	3.9	1.7	0.4	Iris-setosa
6	7	4.6	3.4	1.4	0.3	Iris-setosa
7	8	5.0	3.4	1.5	0.2	Iris-setosa
8	9	4.4	2.9	1.4	0.2	Iris-setosa
9	10	4.9	3.1	1.5	0.1	Iris-setosa

```
In [7]: df.shape #returns the pair (no. of rows , no of columns)
```

```
Out[7]: (150, 6)
```

```
In [8]: df.describe()
```

Out[8]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000

```
In [9]: df.isnull().values.any() #returns true for null values else returns false
```

Out[9]: False

```
In [10]: df.info() # returns the information about the dataset like no of rows n columns
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id               150 non-null   int64
1   SepalLengthCm    150 non-null   float64
2   SepalWidthCm     150 non-null   float64
3   PetalLengthCm    150 non-null   float64
4   PetalWidthCm     150 non-null   float64
5   Species          150 non-null   object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

```
In [11]: df.isnull().sum() # returns no. of missing entries there in in each column
```

```
Out[11]: Id                0
SepalLengthCm            0
SepalWidthCm             0
PetalLengthCm            0
PetalWidthCm             0
Species                  0
dtype: int64
```

```
In [12]: df.head()
```

Out[12]:

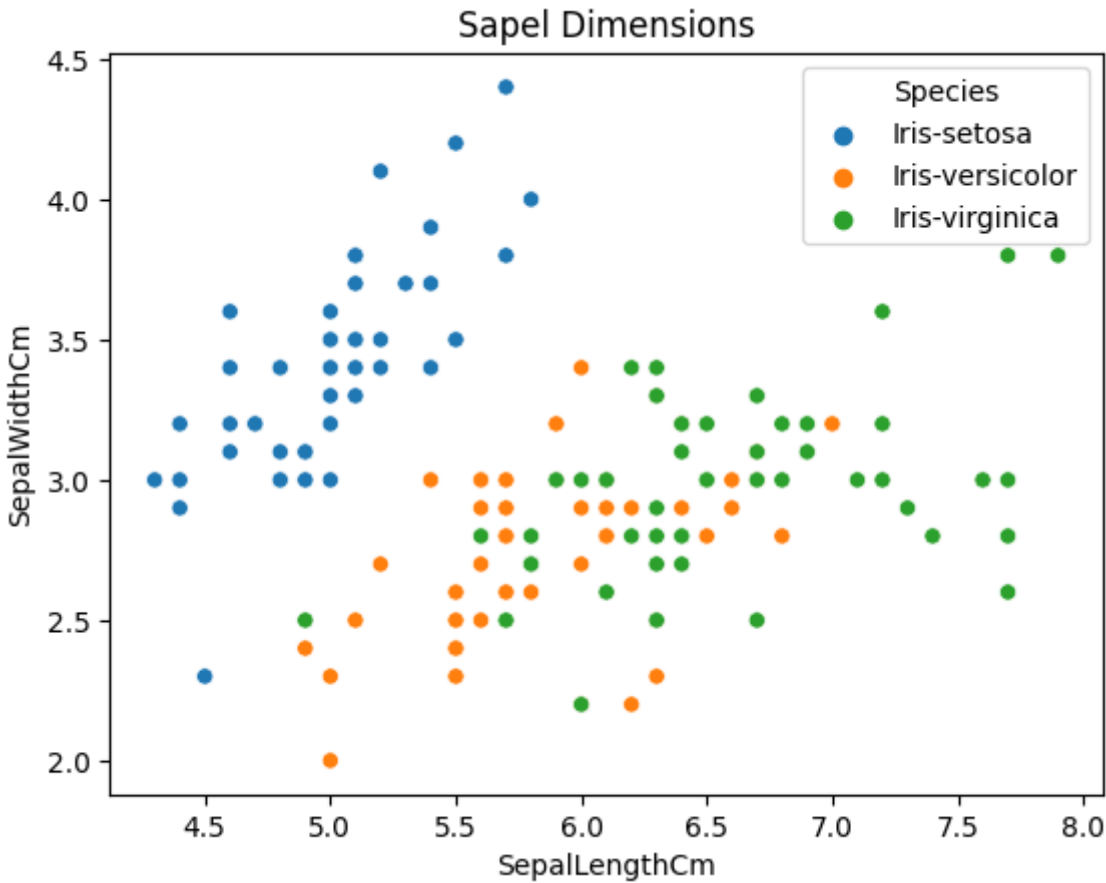
	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

In [13]:

```
plt.title('Sapel Dimensions') #title of plot
sns.scatterplot(x = 'SepalLengthCm' , #seaborn is a visualization library for
                y = 'SepalWidthCm' , #scatterplot() plots a scatter plot grap
                hue = 'Species' ,
                data = df)
```

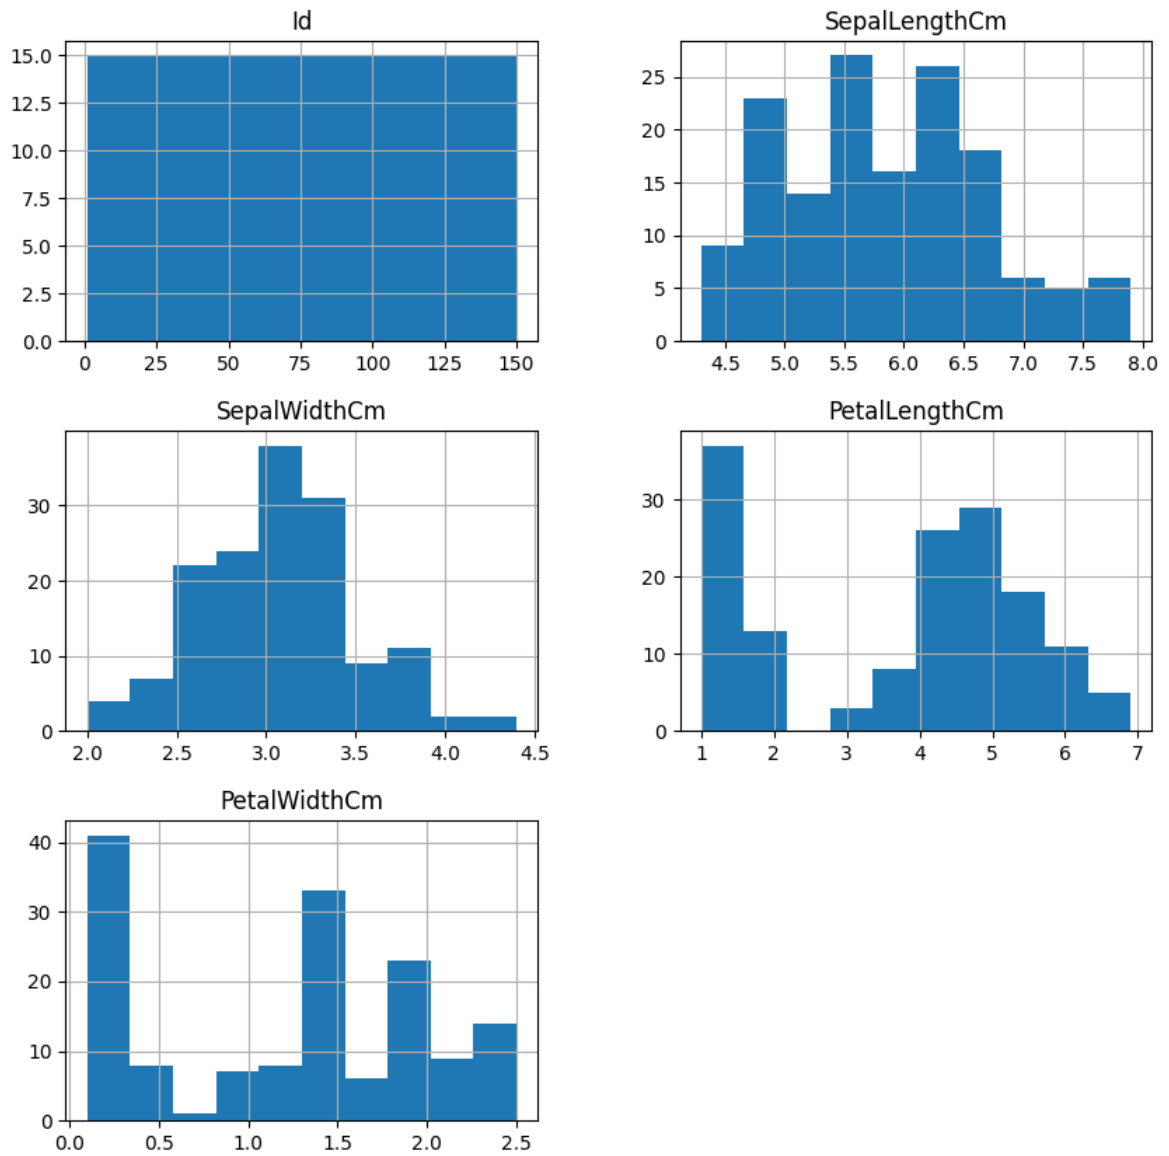
Out[13]:

<Axes: title={'center': 'Sapel Dimensions'}, xlabel='SepalLengthCm', ylabel='SepalWidthCm'>



In [14]:

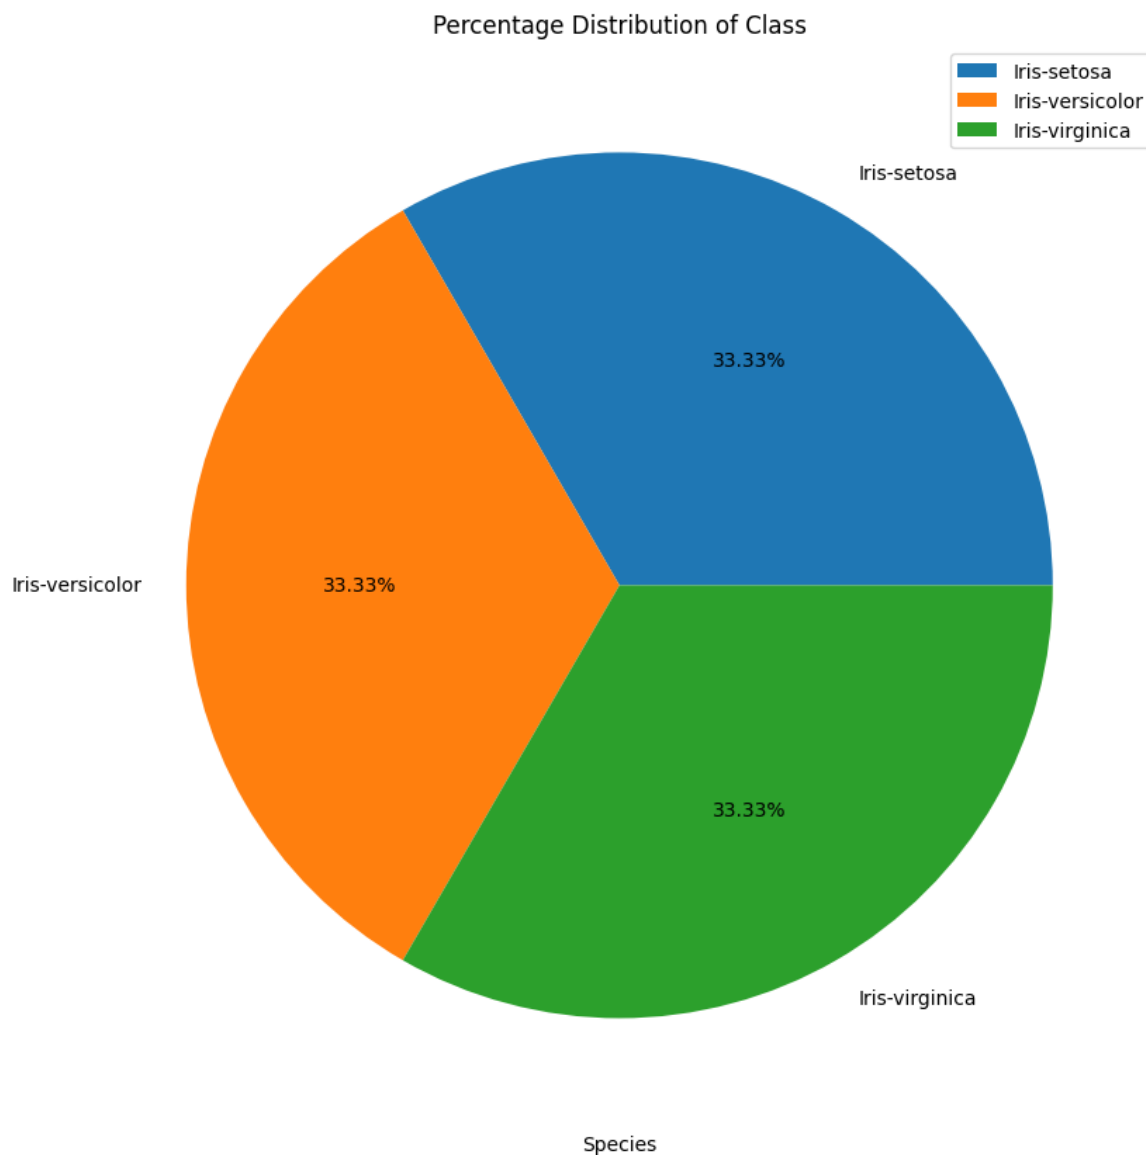
```
df.hist(figsize=(10, 10))
plt.show()
```



```
In [15]: print("\nClass distribution:")
print(df['Species'].value_counts())
```

```
Class distribution:
Species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
Name: count, dtype: int64
```



```
In [16]: # Plot a pie chart to show the percentage distribution of classes
plt.figure(figsize=(20, 10))
df['Species'].value_counts().plot(kind='pie', autopct='%0.2f%%')
plt.title('Percentage Distribution of Class')
plt.legend(df['Species'].value_counts().index)
plt.xlabel('Species')
plt.ylabel(None)
plt.show()
```





```
In [17]: df.columns
```

```
Out[17]: Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',  
              'Species'],  
              dtype='object')
```

```
In [18]: # Show boxplots for each numerical column against the 'class' column  
sns.boxplot(data=df, x='SepalLengthCm', y='Species')  
plt.show()  
sns.boxplot(data=df, x='SepalWidthCm', y='Species')  
plt.show()  
sns.boxplot(data=df, x='PetalLengthCm', y='Species')  
plt.show()  
sns.boxplot(data=df, x='PetalWidthCm', y='Species')  
plt.show()
```









```
In [21]: df.columns
```

```
Out[21]: Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',
               'Species'],
              dtype='object')
```

```
In [22]: # Show violin plots to visualize the distribution of sepal_length for each class
g = sns.violinplot(y='Species', x='SepalLengthCm', data=df, inner='quartile')
plt.show()
g = sns.violinplot(y='Species', x='SepalWidthCm', data=df, inner='quartile')
plt.show()
g = sns.violinplot(y='Species', x='PetalLengthCm', data=df, inner='quartile')
```

```
plt.show()  
g = sns.violinplot(y='Species', x='PetalWidthCm', data=df, inner='quartile')  
plt.show()
```






```
In [23]: df.columns
```

```
Out[23]: Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',  
              'Species'],  
              dtype='object')
```

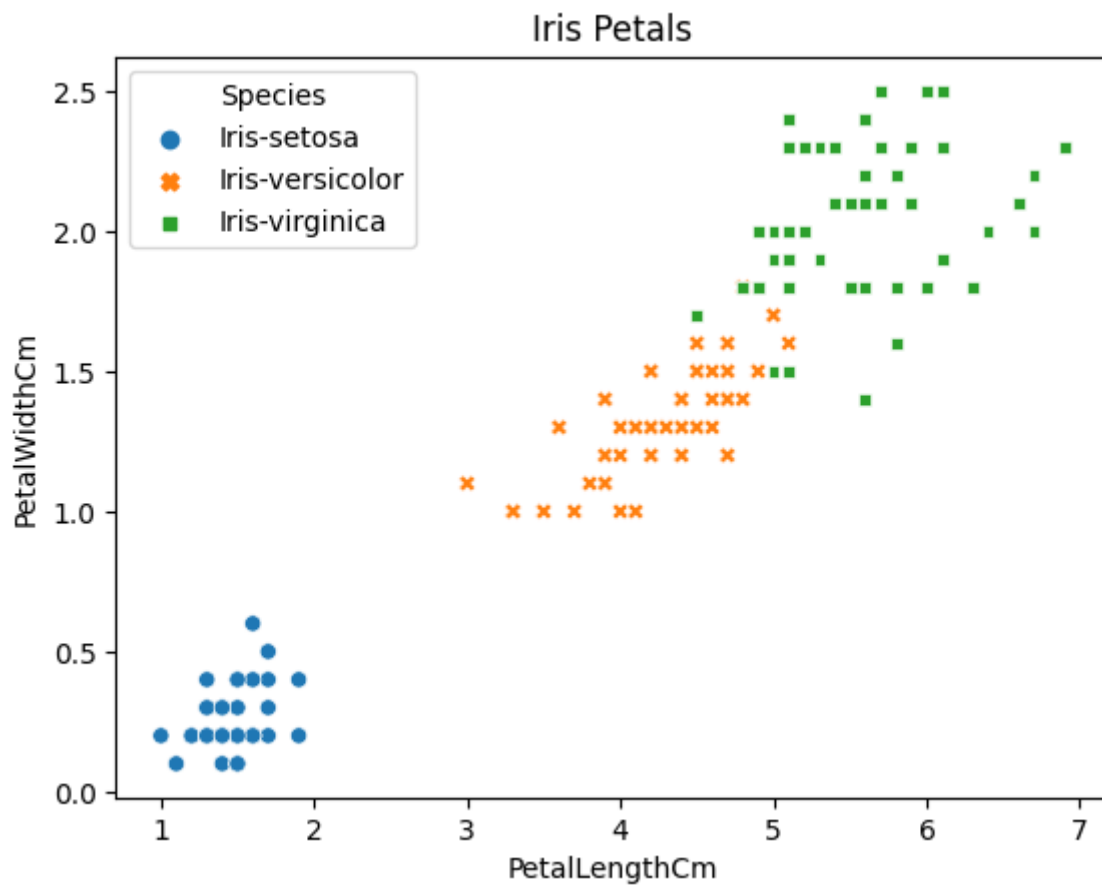
```
In [24]: # Show scatter plots to visualize the relationship between sepal_length and sepal_width  
sns.scatterplot(data=df, x='SepalLengthCm', y='SepalWidthCm', hue='Species', style='Species')  
plt.title('Iris Sepals')  
plt.show()
```



```
In [25]: df.columns
```

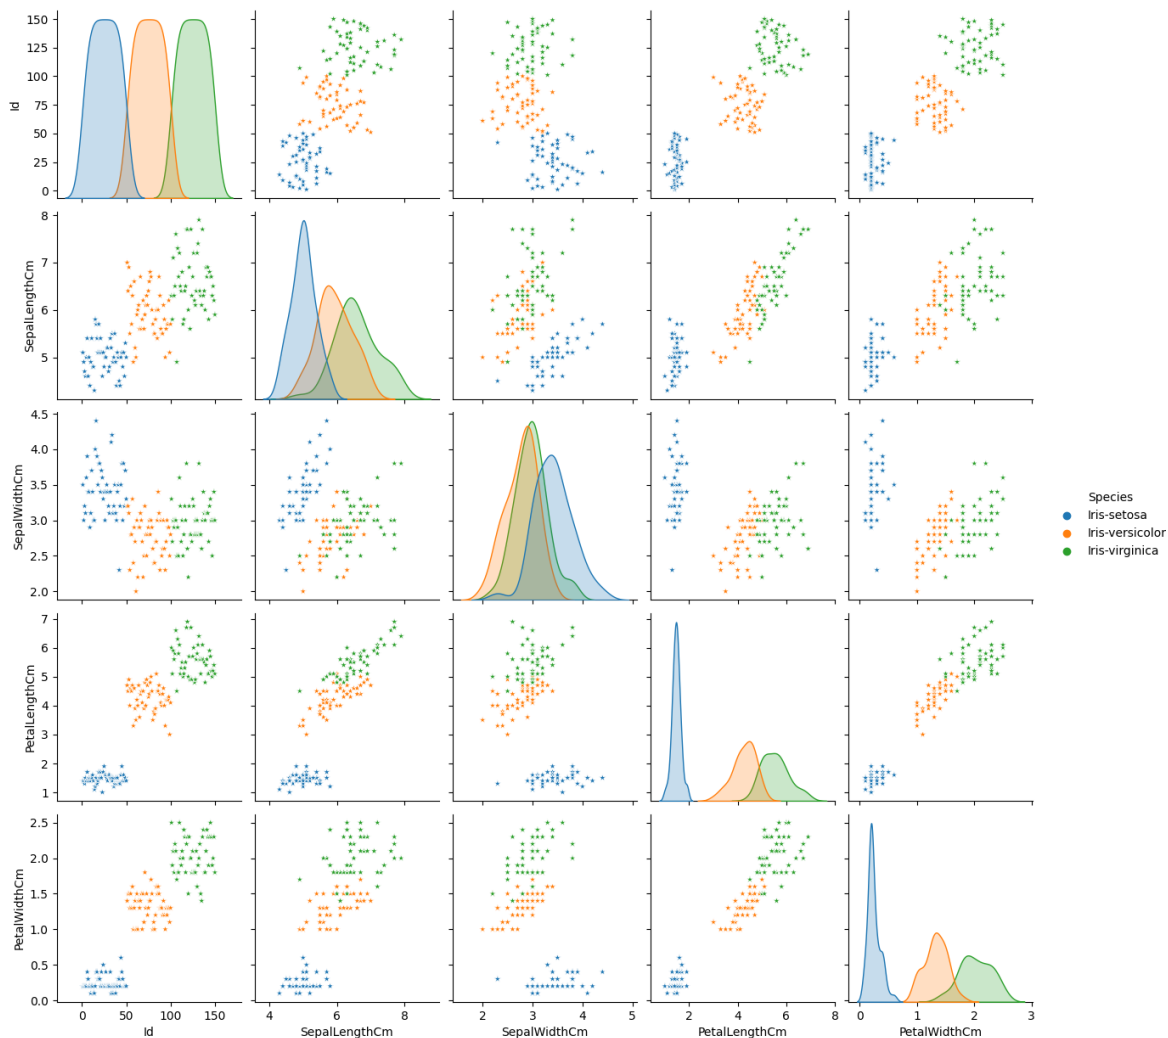
```
Out[25]: Index(['Id', 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm',  
              'Species'],  
             dtype='object')
```

```
In [26]: # Show scatter plots to visualize the relationship between petal_length and petal_width  
sns.scatterplot(data=df, x='PetalLengthCm', y='PetalWidthCm', hue='Species', style='Species')  
plt.title('Iris Petals')  
plt.show()
```



```
In [27]: # Show pair plot to visualize the relationships between all numerical columns with  
sns.pairplot(df, hue='Species', markers='*')
```

```
Out[27]: <seaborn.axisgrid.PairGrid at 0x28955318350>
```



```
In [30]: # Data Encoding
# %matplotlib inline
sns.set_palette('Set1')
```

```
In [31]: # Encode the 'class' column using LabelEncoder
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['Species'] = le.fit_transform(df['Species'])
```

```
In [32]: # Separate features and target variable
x = df.drop(['Species'], axis=1)
y = df['Species']
```

```
In [33]: # Split the data into training and testing sets
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_
```

```
In [34]: # Model Training and Evaluation
from sklearn.metrics import classification_report, confusion_matrix, accuracy_sc
```

```
In [35]: # Support Vector Machine (SVM) classifier
from sklearn.svm import SVC
classifier = SVC()
classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)
print("\nSVM Classifier:")
print(classification_report(y_test, y_pred))
```

```
print(confusion_matrix(y_test, y_pred))
print('Accuracy:', accuracy_score(y_pred, y_test))
```

SVM Classifier:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	7
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0 13  0]
 [ 0  0  7]]
```

Accuracy: 1.0

```
In [36]: # Gaussian Naive Bayes classifier
from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)
print("\nGaussian Naive Bayes Classifier:")
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print('Accuracy:', accuracy_score(y_pred, y_test))
```

Gaussian Naive Bayes Classifier:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	13
2	1.00	1.00	1.00	7
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0 13  0]
 [ 0  0  7]]
```

Accuracy: 1.0

```
In [37]: # Multinomial Naive Bayes classifier
from sklearn.naive_bayes import MultinomialNB
classifier = MultinomialNB()
classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)
print("\nMultinomial Naive Bayes Classifier:")
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print('Accuracy:', accuracy_score(y_pred, y_test))
```

Multinomial Naive Bayes Classifier:

	precision	recall	f1-score	support
0	1.00	0.90	0.95	10
1	0.75	0.69	0.72	13
2	0.56	0.71	0.63	7
accuracy			0.77	30
macro avg	0.77	0.77	0.76	30
weighted avg	0.79	0.77	0.77	30

```
[[9 1 0]
 [0 9 4]
 [0 2 5]]
```

Accuracy: 0.7666666666666667

```
In [38]: # Bernoulli Naive Bayes classifier
from sklearn.naive_bayes import BernoulliNB
classifier = BernoulliNB()
classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)
print("\nBernoulli Naive Bayes Classifier:")
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print('Accuracy:', accuracy_score(y_pred, y_test))
```

Bernoulli Naive Bayes Classifier:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	10
1	0.00	0.00	0.00	13
2	0.23	1.00	0.38	7
accuracy			0.23	30
macro avg	0.08	0.33	0.13	30
weighted avg	0.05	0.23	0.09	30

```
[[ 0  0 10]
 [ 0  0 13]
 [ 0  0  7]]
```

Accuracy: 0.23333333333333334

```
C:\Users\kushw\AppData\Roaming\Python\Python311\site-packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\kushw\AppData\Roaming\Python\Python311\site-packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\kushw\AppData\Roaming\Python\Python311\site-packages\sklearn\metrics\_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

```
In [39]: # Complement Naive Bayes classifier
from sklearn.naive_bayes import ComplementNB
classifier = ComplementNB()
```

```

classifier.fit(x_train, y_train)
y_pred = classifier.predict(x_test)
print("\nComplement Naive Bayes Classifier:")
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print('Accuracy:', accuracy_score(y_pred, y_test))

```

Complement Naive Bayes Classifier:

	precision	recall	f1-score	support
0	0.71	1.00	0.83	10
1	0.00	0.00	0.00	13
2	0.44	1.00	0.61	7
accuracy			0.57	30
macro avg	0.38	0.67	0.48	30
weighted avg	0.34	0.57	0.42	30

```

[[10  0  0]
 [ 4  0  9]
 [ 0  0  7]]

```

Accuracy: 0.5666666666666667

C:\Users\kushw\AppData\Roaming\Python\Python311\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\kushw\AppData\Roaming\Python\Python311\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\kushw\AppData\Roaming\Python\Python311\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```

In [ ]: for clf in classifiers:
        clf.fit(x_train, x_train['sepal_length'])
        name = clf.__class__.__name__

        train_predictions = clf.predict(x_test)
        acc = accuracy_score(x_test['sepal_length'], train_predictions)

        log_entry = pd.DataFrame([[name, acc * 100, 11]], columns=log_cols)
        log = log.append(log_entry)

# Visualize the accuracy of different classifiers using a bar plot
sns.set_color_codes("muted")
sns.barplot(x='Accuracy', y='Classifier', data=log, color="b")
plt.xlabel('Accuracy %')
plt.title('Classifier Accuracy')
plt.show()

# Pie chart to show the distribution of 'sepal_length'
plt.figure(figsize=(8, 8))
sepal_lengths = df['sepal_length']
unique_values = sepal_lengths.unique()
value_counts = sepal_lengths.value_counts()

```

```
plt.pie(value_counts, labels=unique_values, autopct='%.2f%%', startangle=90)
plt.title('Distribution of Sepal Length')
plt.show()
```

```
In [45]: Text = '''In this project, we performed data analysis and classification on the Iris dataset. We started by loading the dataset and performing data visualization to gain insights into the distribution and relationships of the features. We used scatter plots, box plots, violin plots, and pair plots to visualize the relationships between the features and the target classes. These visualizations helped us understand the characteristics of each species and detect any potential outliers. After the data analysis, we encoded the target variable 'Species' using LabelEncoder to convert the categorical class labels into numerical format. Next, we split the data into training and testing sets and trained several classifiers on the training data. We evaluated the performance of each classifier using metrics such as precision, recall, F1-score, and accuracy. The classifiers we used were Support Vector Machine (SVM), Gaussian Naive Bayes, Multinomial Naive Bayes, Bernoulli Naive Bayes, and Complement Naive Bayes. The results of the classification showed that different classifiers achieved varying levels of accuracy and performance on the Iris dataset. The SVM classifier performed well, achieving high accuracy in predicting the species. Gaussian Naive Bayes also showed good performance on this dataset, demonstrating the usefulness of probabilistic models for classification tasks. In conclusion, the Iris dataset is a classic and well-known dataset that serves as a great starting point for exploring data analysis and machine learning classification techniques. By applying various classifiers and visualizing the data, we gained valuable insights into the relationships between features and the classes, and we successfully predicted the species of iris flowers based on their sepal and petal dimensions.'''
print(Text)
```

In this project, we performed data analysis and classification on the Iris dataset using various machine learning models. The Iris dataset contains samples of iris flowers, each belonging to one of three species: Setosa, Versicolor, or Virginica. Our goal was to classify the flowers into their respective species based on their sepal and petal dimensions.

We started by loading the dataset and performing data visualization to gain insights into the distribution and relationships of the features. We used scatter plots, box plots, violin plots, and pair plots to visualize the relationships between the features and the target classes. These visualizations helped us understand the characteristics of each species and detect any potential outliers.

After the data analysis, we encoded the target variable 'Species' using LabelEncoder to convert the categorical class labels into numerical format.

Next, we split the data into training and testing sets and trained several classifiers on the training data. We evaluated the performance of each classifier using metrics such as precision, recall, F1-score, and accuracy. The classifiers we used were Support Vector Machine (SVM), Gaussian Naive Bayes, Multinomial Naive Bayes, Bernoulli Naive Bayes, and Complement Naive Bayes.

The results of the classification showed that different classifiers achieved varying levels of accuracy and performance on the Iris dataset. The SVM classifier performed well, achieving high accuracy in predicting the species. Gaussian Naive Bayes also showed good performance on this dataset, demonstrating the usefulness of probabilistic models for classification tasks.

In conclusion, the Iris dataset is a classic and well-known dataset that serves as a great starting point for exploring data analysis and machine learning classification techniques. By applying various classifiers and visualizing the data, we gained valuable insights into the relationships between features and the classes, and we successfully predicted the species of iris flowers based on their sepal and petal dimensions.

```
In [ ]:
```